

Guide to familial

Ryan Thompson

Introduction

familial is an R package for testing familial hypotheses as described at [insert link to paper]. Briefly, familial hypotheses are statements of the form

$$\begin{aligned} H_0 &: \mu(\lambda) = \mu_0 \text{ for some } \lambda \in \Lambda \\ H_1 &: \mu(\lambda) \neq \mu_0 \text{ for all } \lambda \in \Lambda, \end{aligned}$$

where $\{\mu(\lambda) : \lambda \in \Lambda\}$ is a family of centers. Presently, **familial** supports tests of the Huber family of centers. The mean and median are members of this family.

To test familial hypotheses, **familial** uses the Bayesian bootstrap. The Bayesian bootstrap uses weights $w_1^{(b)}, \dots, w_n^{(b)}$ ($b = 1, \dots, B$) from a uniform Dirichlet distribution in the computation

$$\hat{\mu}^{(b)}(\lambda) := \arg \min_{\mu \in \mathbb{R}} \sum_{i=1}^n w_i^{(b)} l_{\lambda} \left(\frac{x_i - \mu}{\sigma} \right),$$

where the Huber loss function l_{λ} is defined as

$$l_{\lambda}(z) := \begin{cases} \frac{1}{2}z^2 & \text{if } |z| < \lambda \\ \lambda|z| - \frac{1}{2}\lambda^2 & \text{otherwise.} \end{cases}$$

The above minimization is solved for all values of $\lambda \in \Lambda = (0, \infty)$ for $b = 1, \dots, B$. The proportion of times that $\{\hat{\mu}^{(b)}(\lambda) : \lambda \in \Lambda\}_{b=1}^B$ contains the null value μ_0 represents the posterior probability of H_0 being true. To map this probability to a decision (accept, reject, or indeterminate), we assign a loss to making an incorrect decision. The decision that minimizes the expected loss under the posterior distribution is the optimal one.

One-sample tests

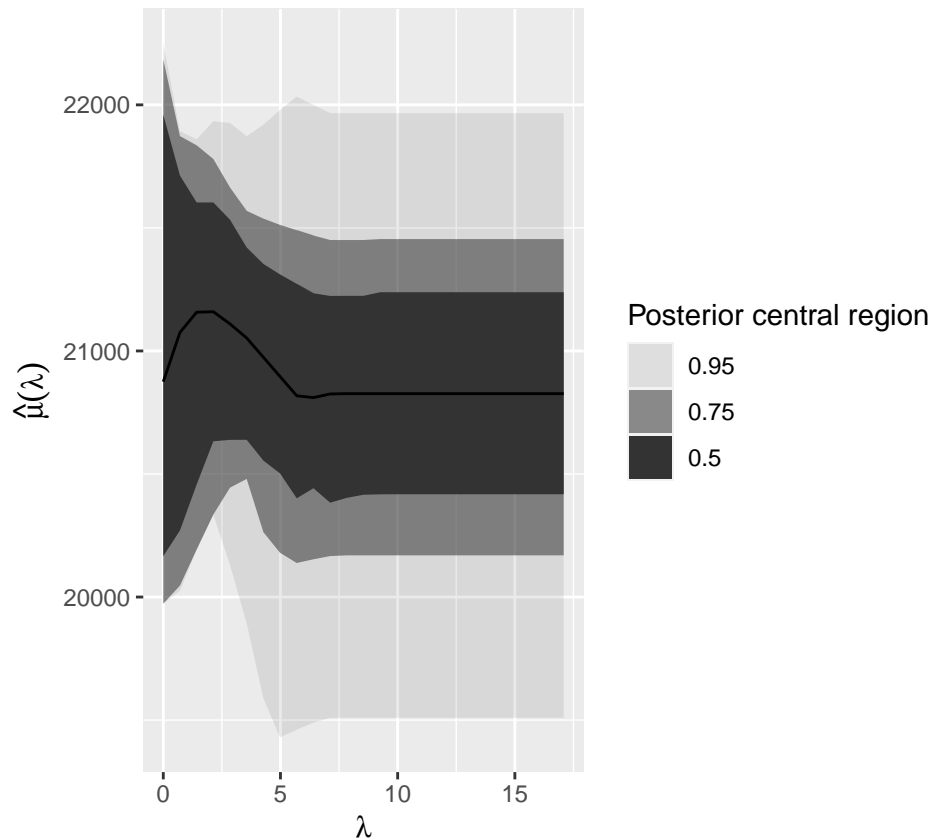
Let's demonstrate the functionality of **familial**. To perform a test of centers with the Huber family, use the **center.test** function with argument **family='huber'** (the default setting). We'll test whether the velocity of galaxies in the **galaxies** dataset is different to 21,000 km/sec.

```
library(familial)
set.seed(1)
x <- MASS::galaxies
test <- center.test(x, mu = 21000)
print(test)
#> -----
#> familial test of centers with huber family
#> -----
#> mu = 21000
#> posterior probabilities:
#>   H0   H1
#> 0.542 0.458
#> optimal decision: indeterminate
```

The output above shows the posterior probabilities for the events H_0 and H_1 . The 54.2% probability assigned to H_0 means that the Huber family contained a center equal to 21,000 km/sec in 54.2% of bootstraps. Because neither of the above probabilities is much larger than the other, the test results in an indeterminate outcome. By default, the function will return an indeterminate result when neither H_0 nor H_1 have probability at least 0.95. This choice of threshold is analogous to using a frequentist significance level of 0.05.

It is possible to visualize the posterior using a functional boxplot via the `plot` function.

```
plot(test)
#> Warning: Using alpha for a discrete variable is not advised.
```



Rather than specify a null value that is a point, we can specify the null as an interval. Let's now test whether the velocity is between 20,500 km/sec and 21,500 km/sec.

```
center.test(x, mu = c(20500, 21500))
#> -----
#> familial test of centers with huber family
#> -----
#> mu = 20500 21500
#> posterior probabilities:
#>   H0   H1
#> 0.959 0.041
#> optimal decision: H0
```

The test now accepts H_0 .

Two-sample tests

`familial` also supports two-sample testing with paired or independent samples. We'll now test whether the weight of cabbages in the `cabbages` dataset is different between two different cultivars. These samples are independent, so we set `paired = FALSE`.


```
x <- MASS::cabbages[MASS::cabbages$Cult == 'c39', 'HeadWt']
y <- MASS::cabbages[MASS::cabbages$Cult == 'c52', 'HeadWt']

test <- center.test(x, y, paired = FALSE)
print(test)
#> -----
#> familial test of centers with huber family
#> -----
#> mu = 0
#> posterior probabilities:
#>   H0   H1
#> 0.008 0.992
#> optimal decision: H1
```

The test rejects H_0 that the weight of cabbages is the same.

We can also visualize the posterior differences between the family of each cultivar via a functional boxplot.

```
plot(test)
#> Warning: Using alpha for a discrete variable is not advised.
```



`familial` also supports one-sided tests. Let's test whether family treatment (FT) improves the weight of anorexia patients in the `anorexia` dataset. These samples are paired.

```

x <- MASS::anorexia[MASS::anorexia$Treat == 'FT', 'Postwt']
y <- MASS::anorexia[MASS::anorexia$Treat == 'FT', 'Prewt']

center.test(x, y, alternative = 'greater', paired = TRUE)
#> -----
#> familial test of centers with huber family
#> -----
#> mu = 0
#> posterior probabilities:
#>     H0    H1
#> 0.006 0.994
#> optimal decision: H1

```

We again reject H_0 that FT does not improve the weight of patients.